



# What is Amazon Bedrock?

[PDF \(/pdfs/bedrock/latest/userguide/bedrock-ug.pdf#what-is-bedrock\)](#)

[RSS \(bedrock-ug.rss\)](#)

Amazon Bedrock is a fully managed service that makes high-performing foundation models (FMs) from leading AI startups and Amazon available for your use through a unified API. You can choose from a wide range of foundation models to find the model that is best suited for your use case. Amazon Bedrock also offers a broad set of capabilities to build generative AI applications with security, privacy, and responsible AI. Using Amazon Bedrock, you can easily experiment with and evaluate top foundation models for your use cases, privately customize them with your data using techniques such as fine-tuning and Retrieval Augmented Generation (RAG), and build agents that execute tasks using your enterprise systems and data sources.

With Amazon Bedrock's serverless experience, you can get started quickly, privately customize foundation models with your own data, and easily and securely integrate and deploy them into your applications using AWS tools without having to manage any infrastructure.

## Topics

- [Features of Amazon Bedrock \(#servicename-feature-overview\)](#)
- [Amazon Bedrock pricing \(#bedrock-pricing\)](#)
- [Supported AWS Regions \(./bedrock-regions.html\)](#)
- [Key definitions \(./key-definitions.html\)](#)



# Features of Amazon Bedrock

Take advantage of Amazon Bedrock foundation models to explore the following capabilities. To see feature limitations by Region, see [Model support by AWS Region \(./models-regions.html\)](#).

- **Experiment with prompts and configurations** – [Run model inference \(./inference.html\)](#) by sending prompts using different configurations and foundation models to generate responses. You can use the API or the text, image, and chat playgrounds in the console to experiment in a graphical interface. When you're ready, set up your application to make requests to the `InvokeModel` APIs.
- **Augment response generation with information from your data sources** – [Create knowledge bases \(./knowledge-base.html\)](#) by uploading data sources to be queried in order to augment a foundation model's generation of responses.
- **Create applications that reason through how to help a customer** – [Build agents \(./agents.html\)](#) that use foundation models, make API calls, and (optionally) query knowledge bases in order to reason through and carry out tasks for your customers.
- **Adapt models to specific tasks and domains with training data** – [Customize an Amazon Bedrock foundation model \(./custom-models.html\)](#) by providing training data for fine-tuning or continued-pretraining in order to adjust a model's parameters and improve its performance on specific tasks or in certain domains.
- **Improve your FM-based application's efficiency and output** – [Purchase Provisioned Throughput \(./prov-throughput.html\)](#) for a foundation model in order to run inference on models more efficiently and at discounted rates.
- **Determine the best model for your use case** – [Evaluate outputs of different models \(./model-evaluation.html\)](#) with built-in or custom prompt datasets to determine the model that is best suited for your application.



**Note**

Model evaluation is in preview release for Amazon Bedrock and is subject to change.

- **Prevent inappropriate or unwanted content** – [Use guardrails](#) ([./guardrails.html](#)) to implement safeguards for your generative AI applications.

## Amazon Bedrock pricing

When you sign up for AWS, your AWS account is automatically signed up for all services in AWS, including Amazon Bedrock. However, you are charged only for the services that you use.

To see your bill, go to the Billing and Cost Management Dashboard in the [AWS Billing and Cost Management console](#) (<https://console.aws.amazon.com/billing/>). To learn more about AWS account billing, see the [AWS Billing User Guide](#) (<https://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/billing-what-is.html>). If you have questions concerning AWS billing and AWS accounts, contact [AWS Support](#) (<https://aws.amazon.com/contact-us/>).

With Amazon Bedrock, you pay to run inference on any of the third-party foundation models. Pricing is based on the volume of input tokens and output tokens, and on whether you have purchased provisioned throughput for the model. For more information, see the [Model providers](#) (<https://console.aws.amazon.com/bedrock/home#/providers>) page in the Amazon Bedrock console. For each model, pricing is listed following the model version. For more information about purchasing Provisioned Throughput, see [Provisioned Throughput for Amazon Bedrock](#) ([./prov-throughput.html](#)).

For more information, see [Amazon Bedrock Pricing](#) (<https://aws.amazon.com/bedrock/pricing>).

**View related pages** Preview





Managedservices › userguide  
[Amazon Bedrock \(AMS  
SSPS\)...](#)

June 27, 2024

AWSCloudFormation › UserGuide  
[AWS::Bedrock::Agent  
InferenceConfiguration...](#)

June 18, 2024

Discover highly rated pages **Preview**



Bedrock › userguide  
[Agents for Amazon  
Bedrock...](#)

July 10, 2024

Bedrock › userguide  
[Amazon Bedrock model  
IDs...](#)

July 23, 2024